

Exercise 03

VU Performance-oriented Computing, Summer Semester 2024

Calvin Hoy

2024-04-07

A) Traditional profiling

Preparation

In order to use `gprof`, `gcc` must be instructed using the flag `-pg` to generate profile information upon running the compiled program, stored in a file named `gmon.out`. To this end, I added `-DCMAKE_C_FLAGS=-pg` to the `cmake` command line.

I also added `-DCMAKE_BUILD_TYPE=RelWithDebInfo`, since profiling an unoptimised build is counterproductive.

Output of `gprof`

Calling `gprof <binary> [<gmon.out>]` without any other arguments prints the flat profile and call graph. Consult the man page for additional information.

I chose to focus on the workload sizes W and B for closer examination.

Note: On LCC3, the module `binutils/2.37` must be loaded, as the version of `gprof` available by default (<2.35) is not able to process the generated profile information.

Flat profile

For brevity, most functions taking less than 1 percent of execution time were omitted from the following output.

W:

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
31.92	0.60	0.60	6712596	0.00	0.00	binvcrhs
17.82	0.94	0.34	6712596	0.00	0.00	matmul_sub
13.03	1.18	0.25	201	1.22	2.94	y_solve
11.70	1.40	0.22	201	1.09	2.82	x_solve
11.17	1.61	0.21	201	1.04	2.77	z_solve
8.51	1.77	0.16	202	0.79	0.79	compute_rhs
3.46	1.84	0.07	6712596	0.00	0.00	matvec_sub
1.06	1.86	0.02	291852	0.00	0.00	binvrhs
1.06	1.88	0.02	291852	0.00	0.00	lhsinit
0.00	1.88	0.00	221472	0.00	0.00	exact_solution
0.00	1.88	0.00	201	0.00	0.00	add
[...]						

The flat profile (captured on my local machine) shows that about 29 percent of execution time is spent in the function `binvcrhs`. A further 15 percent is spent in `matmul_sub`, and just under 14 percent each in `x_solve`/`y_solve`/`z_solve`.

B:

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
28.99	56.57	56.57	609030000	0.00	0.00	binvcrhs
15.11	86.05	29.48	609030000	0.00	0.00	matmul_sub
13.91	113.19	27.14	201	135.02	290.23	y_solve
13.79	140.10	26.91	201	133.88	289.09	x_solve
13.54	166.52	26.42	201	131.46	286.67	z_solve
9.91	185.86	19.33	202	95.71	95.71	compute_rhs
3.53	192.74	6.88	609030000	0.00	0.00	matvec_sub
0.58	193.88	1.14	201	5.67	5.67	add
0.18	194.24	0.36	6030000	0.00	0.00	binvrhs
0.17	194.58	0.34	16980552	0.00	0.00	exact_solution
0.15	194.88	0.30	6030000	0.00	0.00	lhsinit

[...]

For the larger workload size, we see that the number of calls to functions such as `binvcrhs` and `matmul_sub` increases, while e.g. `add` is called equally as often (but its execution time has increased from near-zero to 5.67 ms).

The distribution of time spent does not change for the busiest functions, though e.g. only 0.15 percent of the execution time is spent in `lhsinit` (compared to 1.06 percent for the smaller workload size).

Call graph

index	% time	self	children	called	name
					<spontaneous>
[1]	100.0	0.00	195.12		main [1]
		0.00	193.30	201/201	adi [2]
		1.14	0.00	201/201	add [10]
		0.03	0.26	2/2	initialize [14]
		0.10	0.06	1/1	exact_rhs [15]
		0.00	0.14	1/1	verify [16]
		0.10	0.00	1/1	set_constants [17]
		0.00	0.00	22/22	timer_clear [19]
		0.00	0.00	1/1	timer_start [24]
		0.00	0.00	1/1	timer_stop [25]
		0.00	0.00	1/1	timer_read [23]
		0.00	0.00	1/1	print_results [21]

		0.00	193.30	201/201	main [1]
[2]	99.1	0.00	193.30	201	adi [2]
		27.14	31.20	201/201	y_solve [3]
		26.91	31.20	201/201	x_solve [4]
		26.42	31.20	201/201	z_solve [5]
		19.24	0.00	201/202	compute_rhs [8]

		27.14	31.20	201/201	adi [2]
[3]	29.9	27.14	31.20	201	y_solve [3]
		18.86	0.00	203010000/609030000	binvcrhs [6]
		9.83	0.00	203010000/609030000	matmul_sub [7]
		2.29	0.00	203010000/609030000	matvec_sub [9]
		0.12	0.00	2010000/6030000	binvrhs [11]
		0.10	0.00	2010000/6030000	lhsinit [13]

		26.91	31.20	201/201	adi [2]
[4]	29.8	26.91	31.20	201	x_solve [4]
		18.86	0.00	203010000/609030000	binvcrhs [6]
		9.83	0.00	203010000/609030000	matmul_sub [7]

```

                2.29    0.00 203010000/609030000    matvec_sub [9]
                0.12    0.00 2010000/6030000      binvrhs [11]
                0.10    0.00 2010000/6030000      lhsinit [13]
-----
[5]      29.5    26.42  31.20    201/201          adi [2]
                26.42  31.20    201          z_solve [5]
                18.86    0.00 203010000/609030000    binvrhs [6]
                9.83    0.00 203010000/609030000    matmul_sub [7]
                2.29    0.00 203010000/609030000    matvec_sub [9]
                0.12    0.00 2010000/6030000      binvrhs [11]
                0.10    0.00 2010000/6030000      lhsinit [13]
-----
[6]      29.0    18.86    0.00 203010000/609030000    x_solve [4]
                18.86    0.00 203010000/609030000    y_solve [3]
                18.86    0.00 203010000/609030000    z_solve [5]
                56.57    0.00 609030000          binvrhs [6]
-----
[...]
```

The above (truncated) call graph reveals the following:

- The program spends almost all of its time in the function `adi`.
- `adi` calls `x/y/z_solve` as well as `compute_rhs` and has no computationally intensive code of its own.
- `x/y/z_solve` are *only* called from `adi`.
- `x/y/z_solve` all operate analogously to each other, are called the same number of times, and take up an equal share of execution time.
- `binvrhs` accounts for roughly 19 percent of the time spent in `x/y/z_solve`, and does not call out to any computationally intensive functions.

Annotated source code

`gprof` may also be instructed to output source code with annotations – printing an execution count beside each line with a function definition, as well as a short summary showing which lines (functions) were executed most often.

I found this not to be of much use, as it is essentially just a less concise form of the flat profile.

Comparing with LCC3

The only noteworthy difference I observed in the output of `gprof` on LCC3 is that some or all profile information for the `main` function and other functions called from within it appears to be missing.

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
30.56	96.94	96.94	609030000	0.00	0.00	binvrhs
17.17	151.42	54.48	609030000	0.00	0.00	matmul_sub
13.55	194.42	43.00	201	213.92	488.12	z_solve
11.97	232.39	37.98	201	188.94	463.14	y_solve
11.29	268.21	35.82	201	178.22	452.41	x_solve
10.16	300.45	32.24	202	159.58	159.58	compute_rhs
4.05	313.28	12.83	609030000	0.00	0.00	matvec_sub
0.63	315.28	2.00				add
0.18	315.86	0.58	6030000	0.00	0.00	lhsinit
0.16	316.37	0.51	6030000	0.00	0.00	binvrhs
0.15	316.86	0.49	16980552	0.00	0.00	exact_solution
0.06	317.05	0.19				exact_rhs
0.03	317.15	0.10				set_constants

0.03	317.24	0.09				initialize
0.01	317.26	0.02	1	20.00	50.63	error_norm
0.00	317.27	0.01	1	10.00	10.00	rhs_norm
0.00	317.27	0.00	2	0.00	0.00	wtime_

Call graph

granularity: each sample hit covers 2 byte(s) for 0.00% of 317.27 seconds

index	% time	self	children	called	name
					<spontaneous>
[1]	99.0	0.00	314.21		adi [1]
		43.00	55.11	201/201	z_solve [2]
		37.98	55.11	201/201	y_solve [4]
		35.82	55.11	201/201	x_solve [5]
		32.08	0.00	201/202	compute_rhs [7]

[...]

From the call graph, it appears as though `adi` was called directly at the program entry point. Optimization / inlining?

B) Hybrid trace profiling

Preparation

Tracy must be integrated with the source code to be profiled. Tracy itself is written in C++, and although a C API is provided via `tracy/TracyC.h`, I switched the project language to C++ instead, which necessitated some further adjustments to the code.

The modified source code (including relevant header and source files for Tracy v0.10) may be found in `npb_bt_tracy`. Configure the CMake project with `-DTRACY_ENABLE=On` to create a profiling build.

Frames

Tracy works on the concept of “frames”, as in “frames per second” in e.g. a videogame. The macro `FrameMark` is provided to demark the end of one such frame.

The most natural mapping to me seemed to be the loop over time steps, as it is in effect the main loop of the program:

```
[...]
for (step = 1; step <= niter; step++) {
    if ((step % 20) == 0 || step == 1) {
        printf(" Time step %4d\n", step);
    }

    adi();

    FrameMark;
}
[...]
```

Zone annotations

“Zones” are central to the concept of tracing. Thus, Tracy provides the `ZoneScoped` macro to be added at the start of each interesting function.

Variations on the `ZoneScoped` macro are also available – the `N` suffix for giving the zone a custom name, the `C` suffix for setting a custom colour, and the `S` suffix to capture the call stack at zone entry. None of these were utilised.

Usage of zones is somewhat more involved when using the C API, which is another reason I chose to make the existing source code C++-compatible instead.

I annotated the following functions, most of which were pointed out by `gprof` to have at least 201 executions. The very busiest functions were not annotated so as to limit the number of zones created, and thus the performance impact.

- `initialize`
- `add`
- `adi`
- `exact_rhs`
- `exact_solution`
- `lhsinit`
- `compute_rhs`
- `binvrhs`
- `x/y/z_solve`

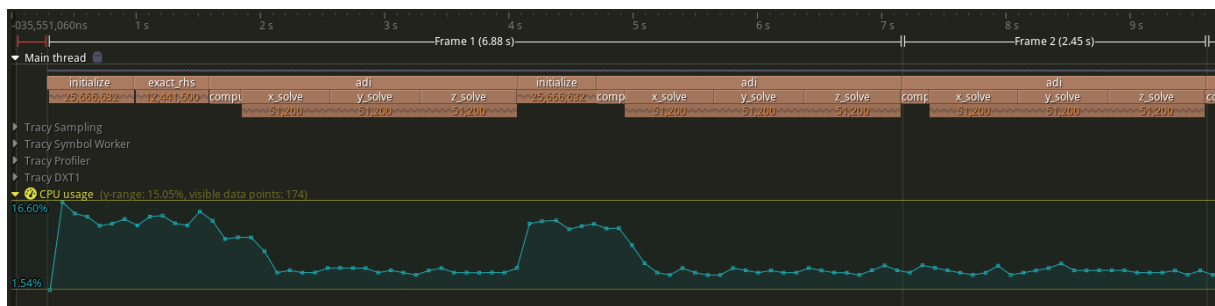


Figure 1: Tracy profiler, workload size C

The section in red before the first frame shows the time that Tracy itself took to initialize.

As I did not put a `FrameMark` between the initialization code and the first run of the loop, the first frame also includes the runtime of all initialization code.

CPU usage is elevated during the two calls to `initialize` – this is due to Tracy itself collecting profiling information for all the zones created by `exact_solution`. For the remaining program duration, CPU usage is reported at around 5 percent – which amounts to full load on a single core on my system.

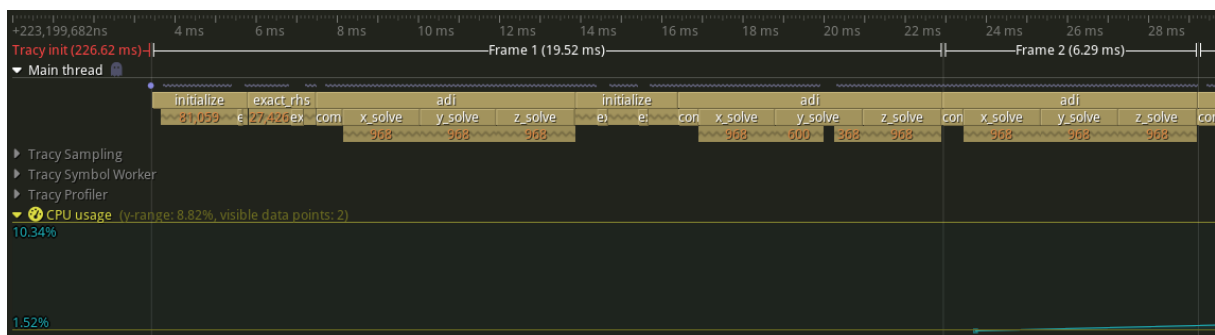


Figure 2: Tracy profiler, workload size W

We can see that the initialization time amounts to roughly 2x the runtime of one iteration of the main loop for both workload sizes W and C.

Measuring performance overhead

I compared the runtime of the modified source code for the workload sizes W, A and C, using the following project configurations:

- none: `-DCMAKE_BUILD_TYPE=RelWithDebInfo`
- tracy: `-DCMAKE_BUILD_TYPE=RelWithDebInfo -DTRACY_ENABLE=On`
- gprof: `-DCMAKE_BUILD_TYPE=RelWithDebInfo -DCMAKE_CXX_FLAGS=-pg`

This amounts in all cases to a build with optimization level `-O2` and debug symbols included.

Note that due to changing the project language to C++, the variable `CMAKE_CXX_FLAGS` was used in place of `CMAKE_C_FLAGS` (as in the first part of this exercise sheet).

All benchmarks were conducted on a Ryzen 9 5900X. Workload size C was tested only once, with all three variants being run simultaneously and bound to cores 0, 2 and 4 of CCD 0 respectively (in order to ensure that none have a thermal / clock speed advantage). The other workload sizes were run in sequence and averaged over five runs each. The Tracy profiling application was not running during the test.

Workload	Configuration	Mean			Variance			Overhead
		wall	user	system	wall	user	system	
W	none	1.206	1.204	0.000	0.001	0.001	0.000	-
	tracy	1.514	1.244	0.158	0.000	0.000	0.000	+25.54%
	gprof	1.252	1.248	0.000	0.001	0.001	0.000	+3.81%
A	none	26.316	26.282	0.028	0.174	0.173	0.000	-
	tracy	29.172	28.188	1.376	0.095	0.099	0.002	+10.85%
	gprof	28.146	28.100	0.038	0.240	0.237	0.000	+6.95%
C	none	461.94	459.90	1.98	-	-	-	-
	tracy	523.89	501.04	22.44	-	-	-	+13.41%
	gprof	499.26	498.26	0.78	-	-	-	+8.08%

Observations

- Tracy, as used here, has a greater performance overhead than `gprof`, and the overhead is most exaggerated for the small workload size.
- The unprofiled program runs almost entirely in userspace (> 99.5 percent). Tracy evidently adds a lot of context switches, raising the time spent in kernel space to just over 4 percent – likely due to its TCP-based communication. `perf stat` (run as root) confirms this:

none:

Performance counter stats for './npb_bt_w':

```
1,223.56 msec task-clock          #    1.000 CPUs utilized
              1      context-switches          #    0.817 /sec
```

[...]

tracy:

Performance counter stats for './npb_bt_w':

```
1,386.54 msec task-clock          #    0.879 CPUs utilized
      1,476      context-switches          #    1.065 K/sec
```

[...]

- Curiously, the `gprof` configuration spent less time in kernel space for workload size C, and this held true for a second run (not included in the table).